

NEURAL NETWORKS

Lecture 14

April 13, 2021



ANNOUNCEMENTS

- Last Lecture Topics
 - Reinforcement Learning
 - Nonlinear dimension reduction (tSNE)
 - NLP (deep learning methods (RNN, LSTM))
- Last Class: April 20
 - Exam 3 Review? Cancel?
- HW 10: Group
- Code 6 is up; Code 7 & 8 this week
- Grade Questions due by Friday April 16




TODAY'S LECTURE

- NN applied to classification
 - Language Model: assigning probabilities to word sequences & predicting upcoming words
- Feedforward network: computation proceeds iteratively from one layer of units to next
- **Deep learning**: modern networks have many layers (~ deep)

HISTORICALLY SPEAKING

- Fundamental tool for language processing
- Derived from McCulloch-Pitts neuron (1943)
 - Human neuron \rightarrow propositional logic computing unit
- Modern NN
 - Network of small computing units
 - Input: vector of values \rightarrow output: one single value

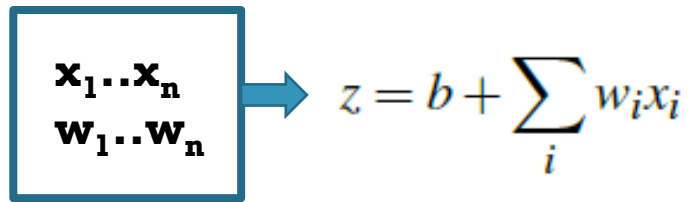


NN VS. LR

- Share similar mathematics with logistic regression
 - More powerful as classifiers
 - Basic NN can learn almost any function
- Different classification approaches
 - LR classifier used on many tasks by developing features based on domain knowledge
 - NN usually take raw words as input & learn features as part of classification process
 - Deep NNs are great for large scale problems with enough training data to automatically learn features

NN UNITS

- NN building block = single computational unit
- Real valued numbers \rightarrow do some computations \rightarrow output
- Computation: weighted sum (z) of input + bias

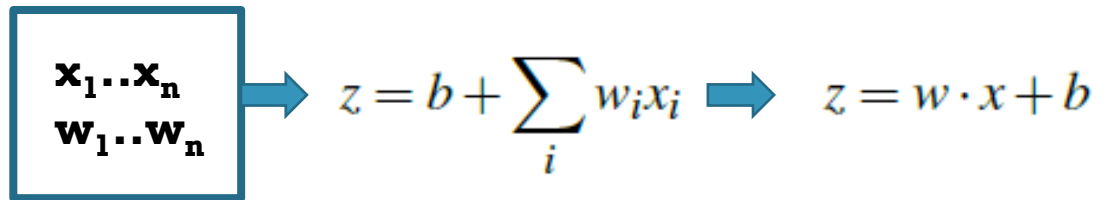


A diagram illustrating the computation within a single computational unit. On the left, a blue-outlined box contains the input variables $x_1 \dots x_n$ and the weight variables $w_1 \dots w_n$. A blue arrow points from this box to the right, where the equation $z = b + \sum_i w_i x_i$ is displayed. The variable z is in blue, b is in orange, and the summation term is in black.

$$\begin{matrix} x_1 \dots x_n \\ w_1 \dots w_n \end{matrix} \rightarrow z = b + \sum_i w_i x_i$$

NN UNITS

- NN building block = single computational unit
- Real valued numbers \rightarrow do some computations \rightarrow output
- Computation: weighted sum (z) of input + bias



The diagram illustrates the computation within a single neural network unit. On the left, a blue-outlined box contains the input vector $\mathbf{x}_1 \dots \mathbf{x}_n$ and the weight vector $\mathbf{w}_1 \dots \mathbf{w}_n$. A blue arrow points from this box to the equation $z = b + \sum_i w_i x_i$. A second blue arrow points from this equation to the simplified dot product form $z = w \cdot x + b$.

$$\begin{array}{c} \mathbf{x}_1 \dots \mathbf{x}_n \\ \mathbf{w}_1 \dots \mathbf{w}_n \end{array} \rightarrow z = b + \sum_i w_i x_i \rightarrow z = w \cdot x + b$$

- Usually expressed as vector notation

ACTIVATION

- Instead of using z (linear function of x) NN units apply non-linear function f to z
- Output of this function == **activation** value for NN unit a
 - If we just model 1 single unit, then the activation for that node/unit is the final output of NN: $y = a = f(z)$

ACTIVATION

- Instead of using z (linear function of x) NN units apply non-linear function f to z
- Output of this function == **activation** value for NN unit a
 - If we just model 1 single unit, then the activation for that node/unit is the final output of NN: $y = a = f(z)$
- 3 popular non-linear (activation) functions $f()$
 - Sigmoid (again!)
 - tanh
 - Rectified linear (ReLU)

SIGMOID

- Maps output into the range $[0, 1]$
 - Good for handling outliers
- Differentiable
 - Good for learning




Figure 7.1 The sigmoid function takes a real value and maps it to the range $[0, 1]$. It is nearly linear around 0 but outlier values get squashed toward 0 or 1.

SIGMOID

- Maps output into the range $[0, 1]$
 - Good for handling outliers
- Differentiable
 - Good for learning

Output of neural unit

$$y = \sigma(w \cdot x + b) = \frac{1}{1 + \exp(-(w \cdot x + b))}$$






Figure 7.1 The sigmoid function takes a real value and maps it to the range $[0, 1]$. It is nearly linear around 0 but outlier values get squashed toward 0 or 1.

SIGMOID


- Input: x_1, x_2, x_3
- Computation
 - Weighted sum ($x_1 w_1 + x_2 w_2 + x_3 w_3$)
 - Adds bias term b
 - Passes sum through sigmoid function
- Output: # between 0..1



SIGMOID




SIGMOID



SIGMOID


$$y = \sigma(w \cdot x + b) = \frac{1}{1 + e^{-(w \cdot x + b)}} = \frac{1}{1 + e^{-(.5 \cdot .2 + .6 \cdot .3 + .1 \cdot .9 + .5)}} = e^{-0.87} = .70$$



ACTIVATION FUNCTIONS

tanh $y = \frac{e^z - e^{-z}}{e^z + e^{-z}}$


- Very similar to sigmoid, but usually better
- Sigmoid variant
- Range: $[-1, 1]$



ACTIVATION FUNCTIONS


tanh $y = \frac{e^z - e^{-z}}{e^z + e^{-z}}$

- Very similar to sigmoid, but usually better
- Sigmoid variant
- Range: $[-1, 1]$




ReLU $y = \max(x, 0)$

- Rectified Linear Unit
- Simplest
- Most commonly used
- Avoids saturation problem



WHY NEURAL NETWORKS?

- Combine neural units into larger & larger networks just like biological neurons
- Minsky & Papert (1969): single neural unit cannot compute simple functions of its input (so need layers)



XOR PROBLEM

AND		
X1	X2	Y
0	0	0
0	1	0
1	0	0
1	1	1

OR		
X1	X2	Y
0	0	0
0	1	1
1	0	1
1	1	1

XOR		
X1	X2	Y
0	0	0
0	1	1
1	0	1
1	1	0

XOR PROBLEM

AND		
X1	X2	Y
0	0	0
0	1	0
1	0	0
1	1	1

OR		
X1	X2	Y
0	0	0
0	1	1
1	0	1
1	1	1

XOR		
X1	X2	Y
0	0	0
0	1	1
1	0	1
1	1	0



- M & P used **perceptron**
- Simple neural unit
- Binary output $y = 0$ or 1
- No non-linear activation function

$$y = \begin{cases} 0, & \text{if } w \cdot x + b \leq 0 \\ 1, & \text{if } w \cdot x + b > 0 \end{cases}$$

XOR PROBLEM

AND		
X1	X2	Y
0	0	0
0	1	0
1	0	0
1	1	1

OR		
X1	X2	Y
0	0	0
0	1	1
1	0	1
1	1	1





XOR PROBLEM

AND		
X1	X2	Y
0	0	0
0	1	0
1	0	0
1	1	1

OR		
X1	X2	Y
0	0	0
0	1	1
1	0	1
1	1	1

XOR		
X1	X2	Y
0	0	0
0	1	1
1	0	1
1	1	0





XOR PROBLEM

AND		
X1	X2	Y
0	0	0
0	1	0
1	0	0
1	1	1

OR		
X1	X2	Y
0	0	0
0	1	1
1	0	1
1	1	1


XOR		
X1	X2	Y
0	0	0
0	1	1
1	0	1
1	1	0

- Perceptron for XOR?
- Linear classifier
 - For 2D input x_1 & x_2
 $w_1x_1 + w_2x_2 + b = 0$
 - It's a line \rightarrow decision boundary




XOR PROBLEM

AND		
X1	X2	Y
0	0	0
0	1	0
1	0	0
1	1	1



a) x_1 AND x_2




b) x_1 OR x_2


OR		
X1	X2	Y
0	0	0
0	1	1
1	0	1
1	1	1

XOR PROBLEM

AND		
X1	X2	Y
0	0	0
0	1	0
1	0	0
1	1	1



a) x_1 AND x_2




b) x_1 OR x_2


OR		
X1	X2	Y
0	0	0
0	1	1
1	0	1
1	1	1

XOR PROBLEM

AND		
X1	X2	Y
0	0	0
0	1	0
1	0	0
1	1	1



a) x_1 AND x_2




b) x_1 OR x_2

OR		
X1	X2	Y
0	0	0
0	1	1
1	0	1
1	1	1

Linearly Separable

XOR PROBLEM

XOR		
X1	X2	Y
0	0	0
0	1	1
1	0	1
1	1	0




c) $x_1 \text{ XOR } x_2$

Not Linearly Separable


SOLUTION: NEURAL NETWORKS!

- Calculate XOR with a layered network of units
- Goodfellow et al. solution: use 2 layers of ReLU-based units




SOLUTION: NEURAL NETWORKS!

- Calculate XOR with a layered network of units
- Goodfellow et al. solution: use 2 layers of ReLU-based units



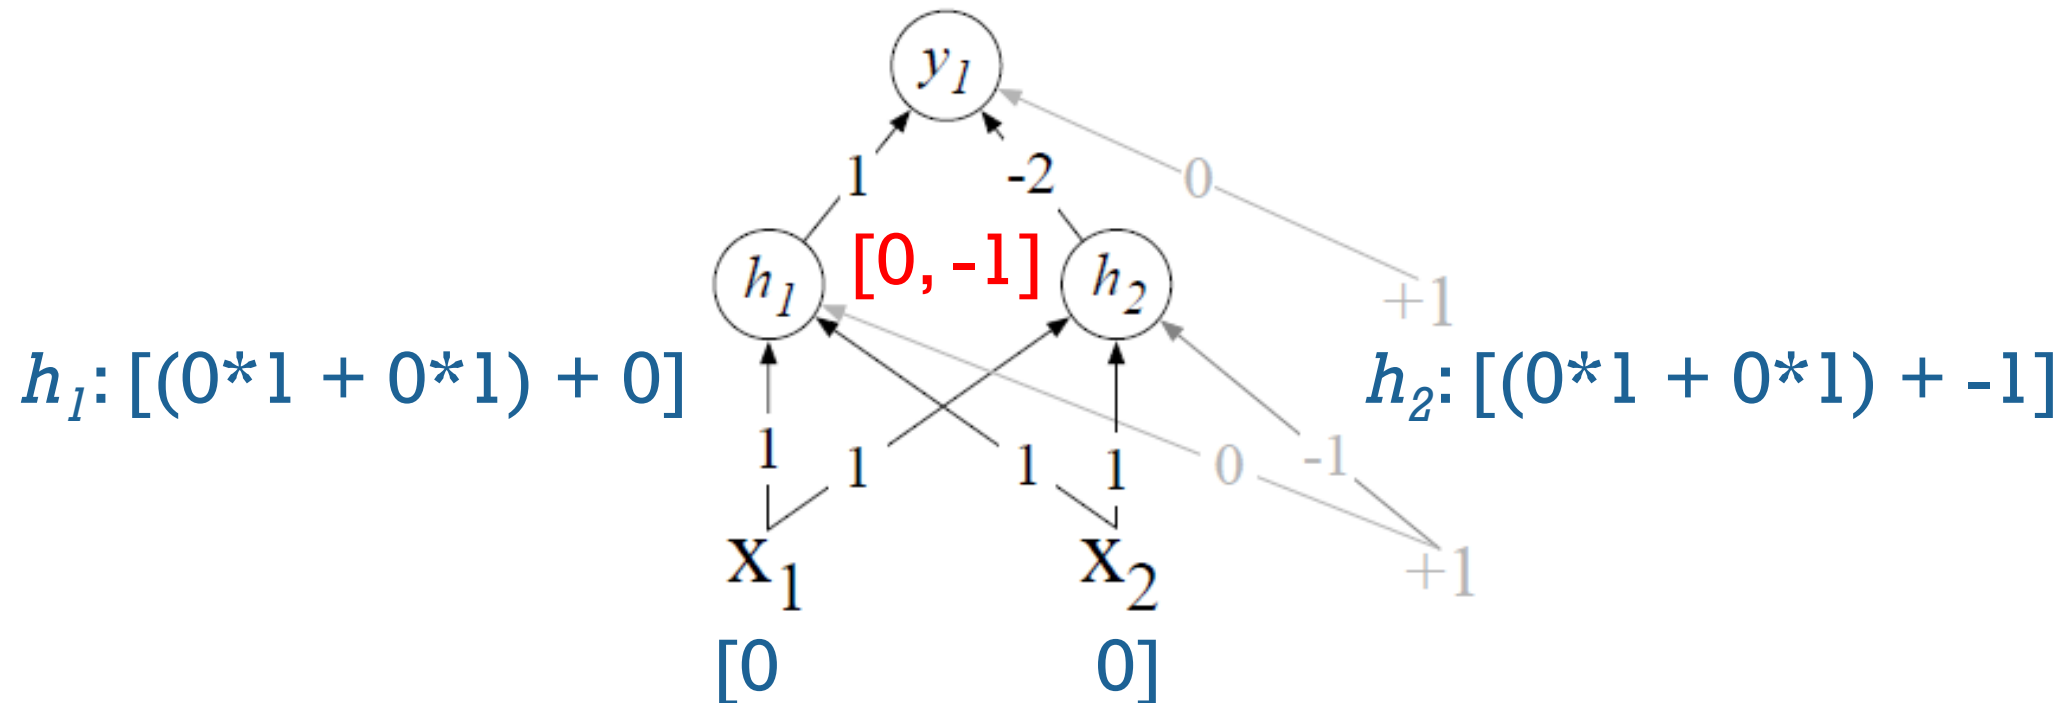
SOLUTION: NEURAL NETWORKS!

- Calculate XOR with a layered network of units
- Goodfellow et al. solution: use 2 layers of ReLU-based units



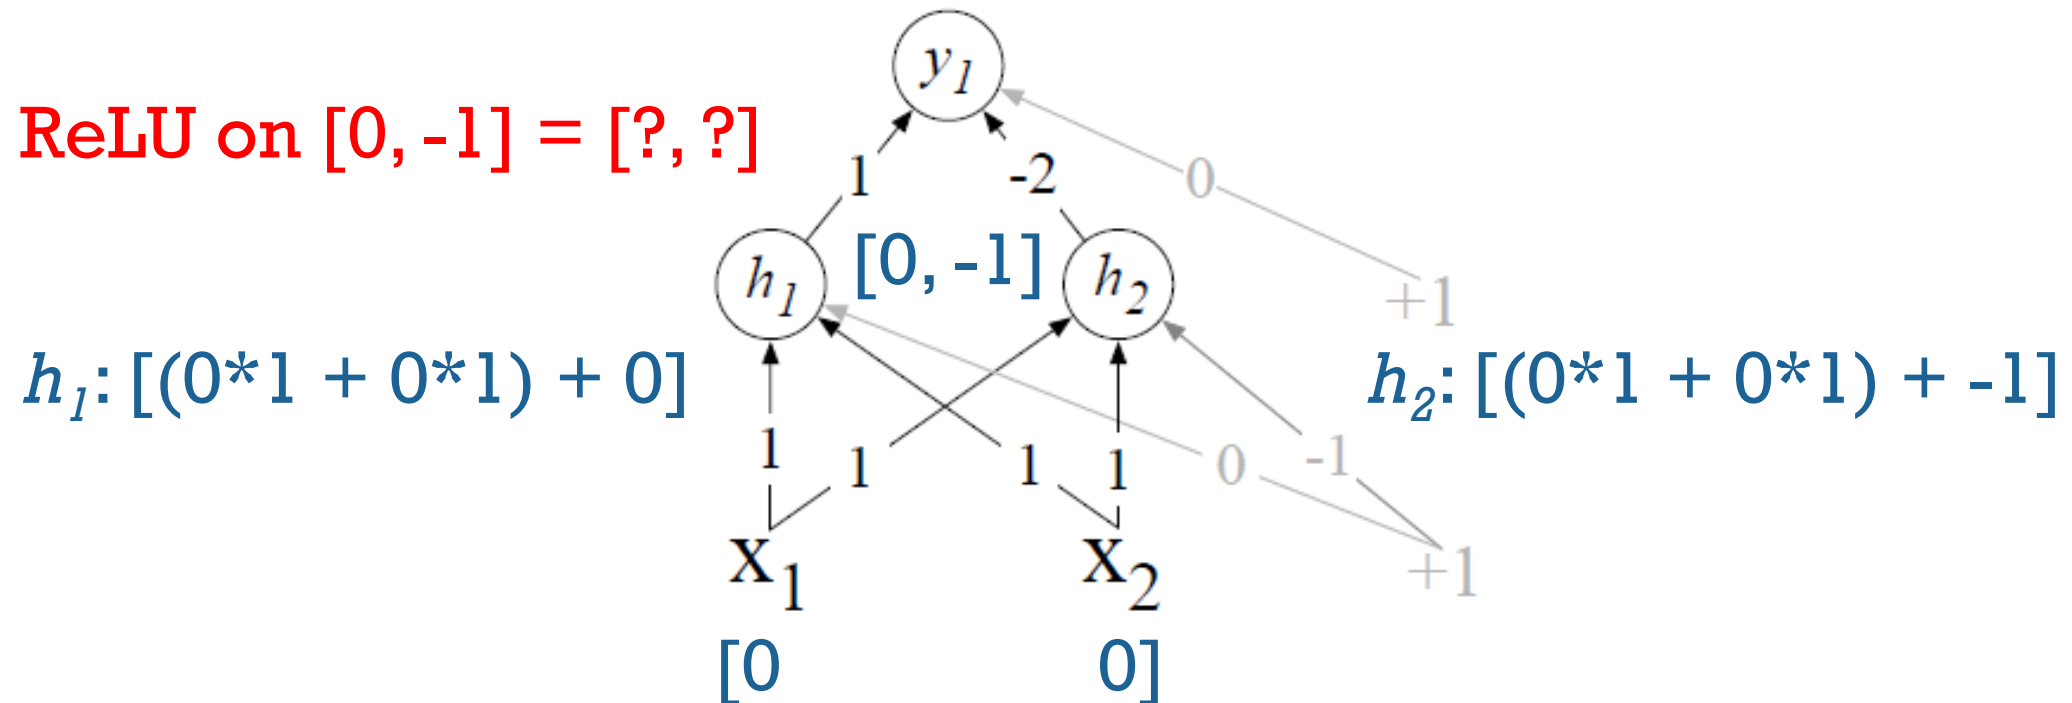
SOLUTION: NEURAL NETWORKS!

- Calculate XOR with a layered network of units
- Goodfellow et al. solution: use 2 layers of ReLU-based units



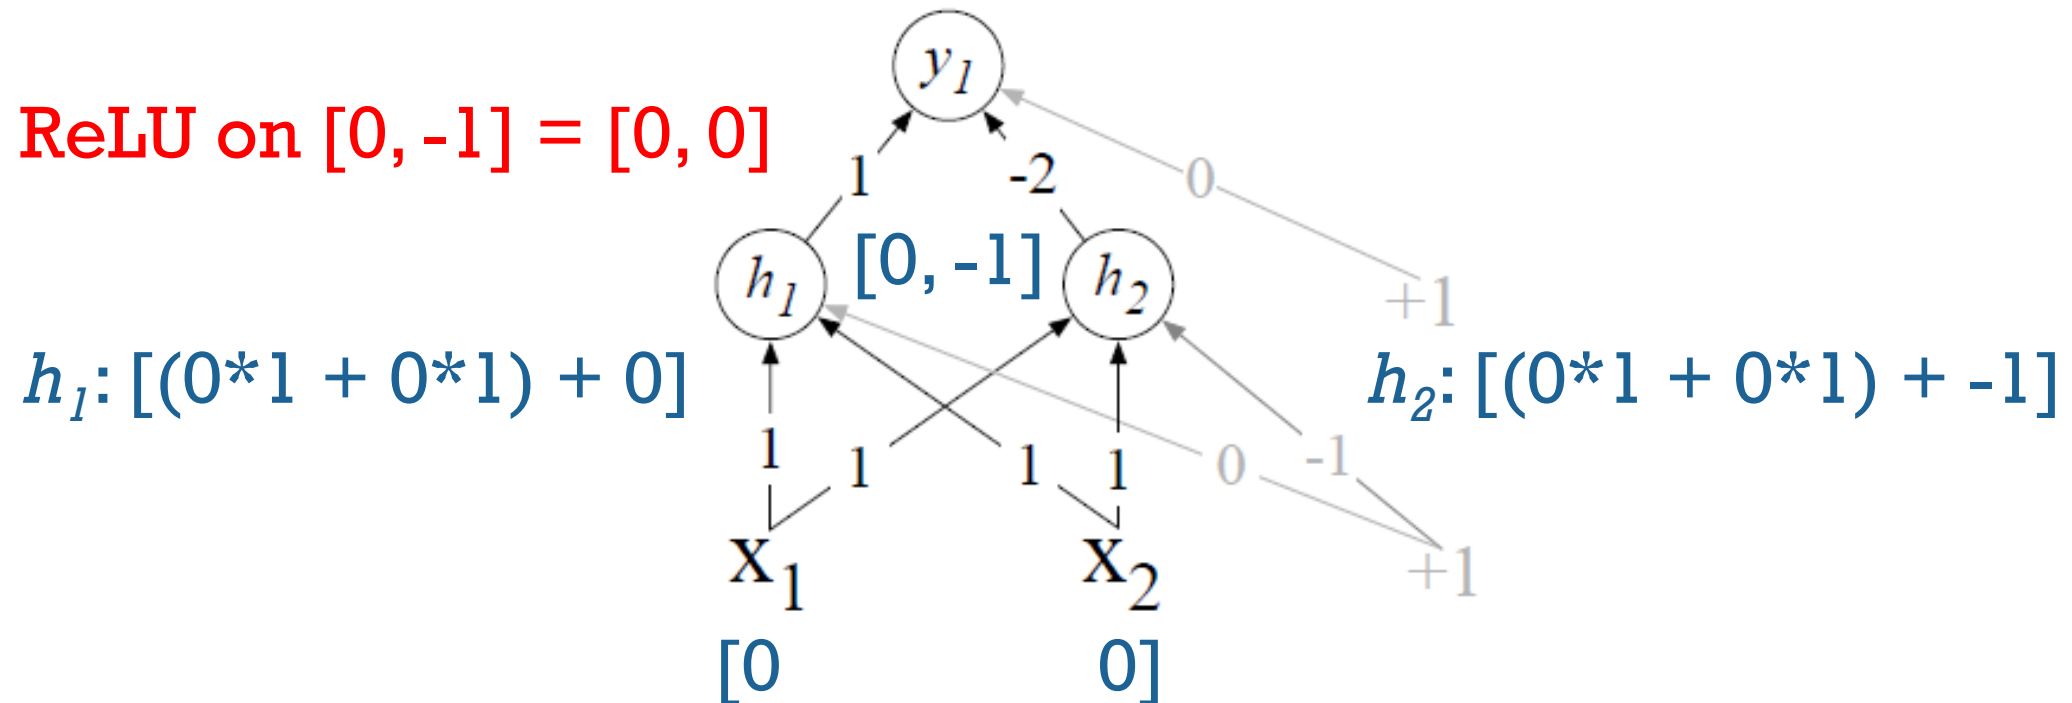
SOLUTION: NEURAL NETWORKS!

- Calculate XOR with a layered network of units
- Goodfellow et al. solution: use 2 layers of ReLU-based units



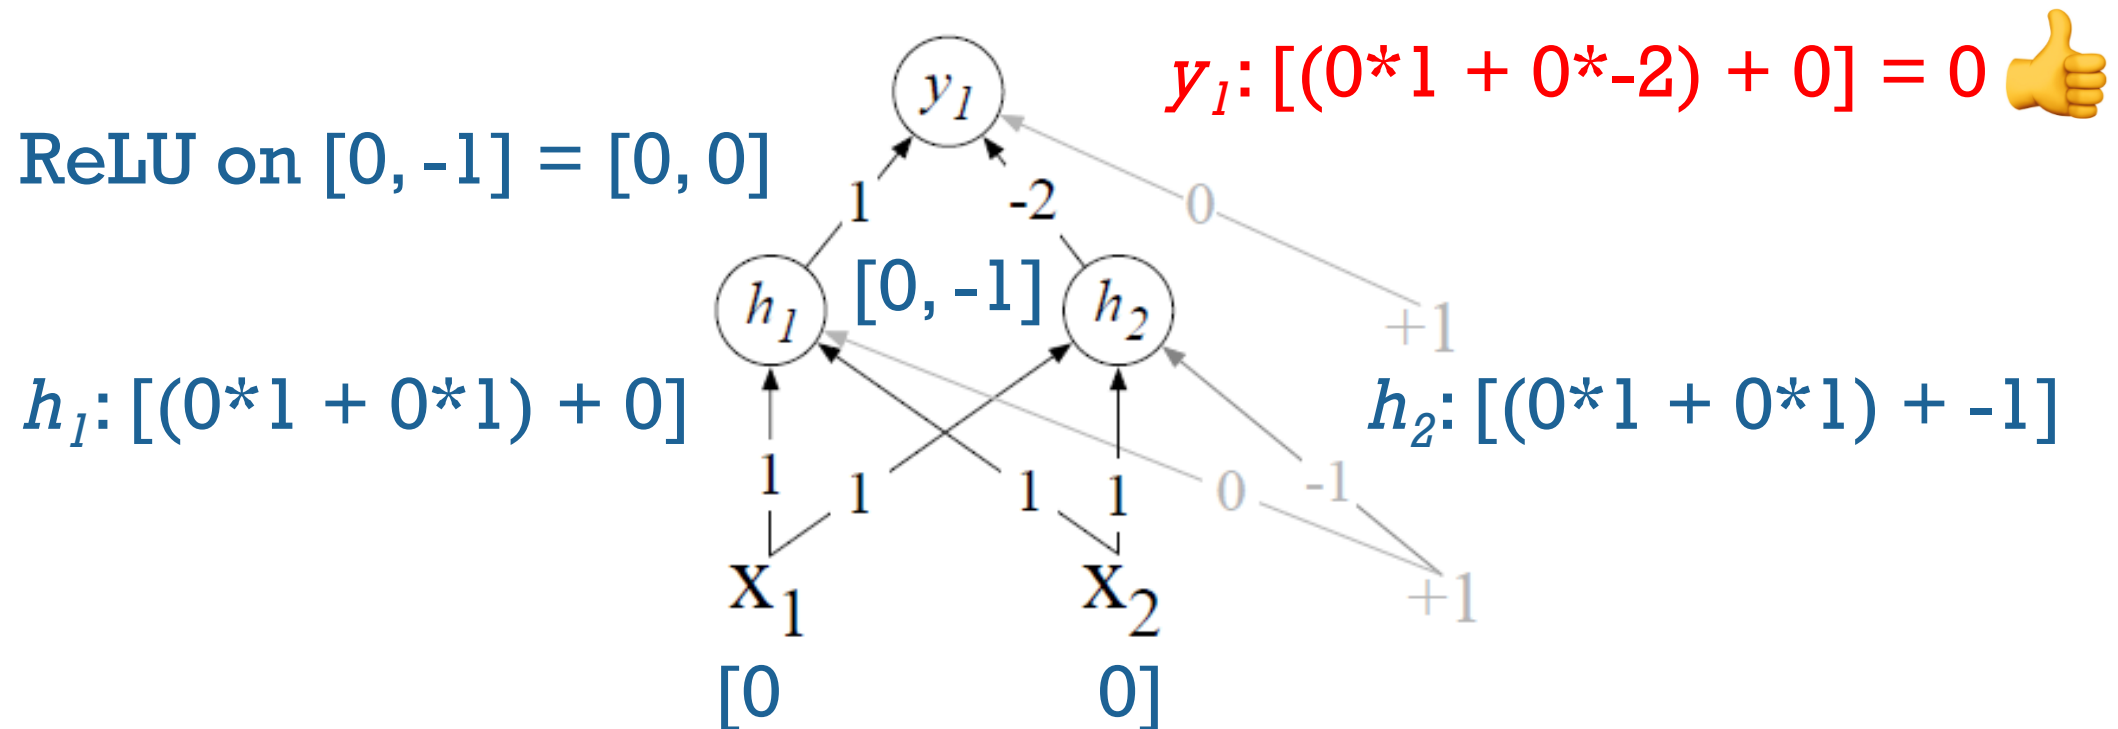
SOLUTION: NEURAL NETWORKS!

- Calculate XOR with a layered network of units
- Goodfellow et al. solution: use 2 layers of ReLU-based units



SOLUTION: NEURAL NETWORKS!

- Calculate XOR with a layered network of units
- Goodfellow et al. solution: use 2 layers of ReLU-based units




SOLUTION: NEURAL NETWORKS!

- Calculate XOR with a layered network of units
- Goodfellow et al. solution: use 2 layers of ReLU-based units

ReLU on $[0, -1] = [0, 0]$

Here we gave the NN the weights but NNs auto-learn their weights through error backpropagation → so NNs can automatically learn useful representations of input



$$y_1: [(0*1 + 0*-2) + 0] = 0 \text{ 👍}$$

$$h_2: [(0*1 + 0*1) + -1]$$

FEED FORWARD NEURAL NETWORKS

- Multi-layer network of connected units with no cycles
 - No cycles = unit outputs from each layer are sent to the next higher layer & no output is passed back to lower layers
- Multi-layer feed forward network $\{=, \neq\}$ multi-layer perceptron (MLP)???

FEED FORWARD NEURAL NETWORKS

- 3 kinds of nodes
 - Input units
 - Hidden units
 - Output units

FEED FORWARD NEURAL NETWORKS


- 3 kinds of nodes
 - Input units: scalar values
 - Hidden units
 - Output units

FEED FORWARD NEURAL NETWORKS

- 3 kinds of nodes
 - Input units: scalar values
 - Hidden units: take weighted sum of inputs & apply activation function (non-linear function)
 - Output units


FEED FORWARD NEURAL NETWORKS

- 3 kinds of nodes
 - Input units: scalar values
 - Hidden units: take weighted sum of inputs & apply activation function (non-linear function)
 - Output units: answer!



FEED FORWARD NEURAL NETWORKS

- **Hidden layer**: core of the NN, composed of hidden units
- **Fully-connected**
 - Input of each unit in each layer is output from all units in the previous layer
 - Link between every pair of units from 2 adjacent layers



HIDDEN LAYER

- Hidden unit parameters: vector for w & b scalar
- Hidden layer parameters: W matrix & \mathbf{b} vector
 - Parameters for entire hidden layer (all hidden units)
 - W combines weight vector w_i for each hidden unit h_i
 - W_{ij} = weight of connection from i^{th} input x_i to j^{th} hidden unit h_j
 - \mathbf{b} combines bias b_i for each hidden unit h_i
- Advantage of using a single matrix W to represent weights for entire layer?

HIDDEN LAYER

- Efficient hidden layer computation using simple matrix operations
 - 1. Multiply weight matrix W by input vector \mathbf{x}
 - 2. Add bias vector \mathbf{b}
 - 3. Apply activation function g (element-wise)

HIDDEN LAYER

- Efficient hidden layer computation using simple matrix operations
 - 1. Multiply weight matrix W by input vector \mathbf{x}
 - 2. Add bias vector \mathbf{b}
 - 3. Apply activation function g (element-wise)

$$h = \sigma(Wx + b)$$

HOW DO WE GET Y ?

- Hidden layer forms a representation h of input
- Output layer takes h & computes final output value

HOW DO WE GET Y?

- Real valued #
- If classification is NN's goal:
 - Binary task (e.g., sentiment classification)
 - y = probability of positive vs. negative sentiment
 - Multinomial classification (e.g., part-of-speech (POS) tagging)
 - 1 output node for each POS where value = probability of that POS
 - all values of output nodes must sum to 1
 - output layer \approx probability distribution across output nodes

HOW DO WE GET Y?

- Output layer has a weight matrix U (bias vector optional)
- Intermediate output $z = \text{weight matrix } U * \text{input vector } h$
- At this point, what's wrong with z ?

$$z = Uh$$

HOW DO WE GET Y ?

- z is a vector of real-valued #s
- For classification need a vector of probabilities

HOW DO WE GET Y?

- z is a vector of real-valued #s
- For classification need a vector of probabilities
- Normalize!
- **Softmax** function “normalizes” a vector of real values by converting it into one that encodes a probability distribution (all #s are between 0..1 & sum to 1)

SOFTMAX

- Just plug in real values from z $\text{softmax}(z_i) = \frac{e^{z_i}}{\sum_{j=1}^d e^{z_j}} \quad 1 \leq i \leq d$
- LR uses to create a probability distribution from sum of weights * features

NN VS. LR


- Think of NN as classifier with 1 hidden layer
 - Build input representation (hidden) vector h
 - Run standard LR on features the NN developed in h

NN VS. LR

- Think of NN as classifier with 1 hidden layer
 - Build input representation (hidden) vector h
 - Run standard LR on features the NN developed in h
- How NN differs from LR
 - Many layers (deep NN \approx layer on layer of LR classifiers)
 - Instead of using feature templates/engineering use previous layers to induce feature representations

FEEDFORWARD NN

$$\begin{aligned}h &= \sigma(Wx + b) \\z &= Uh \\y &= \textit{softmax}(z)\end{aligned}$$



FEEDFORWARD NN

$$h = \sigma(Wx + b)$$

$$z = Uh$$

$$y = \text{softmax}(z)$$


$$z^{[1]} = W^{[1]}a^{[0]} + b^{[1]}$$

$$a^{[1]} = g^{[1]}(z^{[1]})$$

$$z^{[2]} = W^{[2]}a^{[1]} + b^{[2]}$$

$$a^{[2]} = g^{[2]}(z^{[2]})$$

$$\hat{y} = a^{[2]}$$



FEEDFORWARD NN

$$h = \sigma(Wx + b)$$

$$z = Uh$$

$$y = \text{softmax}(z)$$

$$z^{[1]} = W^{[1]}a^{[0]} + b^{[1]}$$

$$a^{[1]} = g^{[1]}(z^{[1]})$$

$$z^{[2]} = W^{[2]}a^{[1]} + b^{[2]}$$

$$a^{[2]} = g^{[2]}(z^{[2]})$$

$$\hat{y} = a^{[2]}$$




for i in 1..n

$$z^{[i]} = W^{[i]}a^{[i-1]} + b^{[i]}$$

$$a^{[i]} = g^{[i]}(z^{[i]})$$

$$\hat{y} = a^{[n]}$$



TRAINING NN

- Feedforward NN is instance of supervised ML
- We know the correct output y for each observation x
- NN produces y^* (prediction; estimate of true y)

TRAINING NN

- Goal of training: learn parameters
 - Learn $W^{[i]}$ & $b^{[i]}$ for each layer i that makes y^* for each training observation as close as possible to true y

TRAINING NN

- Goal of training: learn parameters
 - Learn $W^{[i]}$ & $b^{[i]}$ for each layer i that makes y^* for each training observation as close as possible to true y
- Follow same steps as LR
 - Cross-entropy loss: to model distance between y^* & y
 - Gradient descent: to find parameters to minimize loss
 - Optimization is tricky now!

TRAINING NN

- Optimization: GD requires knowing gradient of loss function
 - Vector contains partial derivative of loss function w.r.t. each parameter
- In LR: directly compute derivative
- In NN: how can we compute the partial derivative of some weight in layer 1 when loss is attached to later layer?

TRAINING NN

- Optimization: GD requires knowing gradient of loss function
 - Vector contains partial derivative of loss function w.r.t. each parameter
- In LR: directly compute derivative
- In NN: how can we compute the partial derivative of some weight in layer 1 when loss is attached to later layer?
 - Error backpropagation or reverse differentiation

TRAINING NN: CROSS-ENTROPY LOSS

- Binary classification (with sigmoid on final output layer)
== LR loss equation

$$L_{CE}(\hat{y}, y) = -\log p(y|x) = -[y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$$

TRAINING NN: CROSS-ENTROPY LOSS

- Binary classification (with sigmoid on final output layer)
== LR loss equation

$$L_{CE}(\hat{y}, y) = -\log p(y|x) = -[y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$$

- Multinomial classification $L_{CE}(\hat{y}, y) = -\sum_{i=1}^C y_i \log \hat{y}_i$

TRAINING NN: CROSS-ENTROPY LOSS

- Binary classification (with sigmoid on final output layer)
== LR loss equation

$$L_{CE}(\hat{y}, y) = -\log p(y|x) = -[y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$$

- Multinomial classification $L_{CE}(\hat{y}, y) = -\sum_{i=1}^C y_i \log \hat{y}_i$

- Hard classification (only 1 class is correct)

$$L_{CE}(\hat{y}, y) = -\log \hat{y}_i$$

negative log likelihood loss

TRAINING NN: CROSS-ENTROPY LOSS

- Binary classification (with sigmoid on final output layer)
== LR loss equation

$$L_{CE}(\hat{y}, y) = -\log p(y|x) = -[y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$$

- Multinomial classification $L_{CE}(\hat{y}, y) = -\sum_{i=1}^C y_i \log \hat{y}_i$

- Hard classification (only 1 class is correct)

$$L_{CE}(\hat{y}, y) = -\log \hat{y}_i$$

negative log likelihood loss

softmax
K classes

$$L_{CE}(\hat{y}, y) = -\log \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$

TRAINING NN: COMPUTE GRADIENT

- Gradient = partial derivative of loss function w.r.t. each parameter
- Simple cases
 - 1 weight layer & sigmoid output → derivative of LR loss
 - 1 hidden layer & softmax output → derivative of softmax

TRAINING NN: COMPUTE GRADIENT

- Gradient = partial derivative of loss function w.r.t. each parameter
- Simple cases
 - 1 weight layer & sigmoid output → derivative of LR loss
 - 1 hidden layer & softmax output → derivative of softmax
- Problems
 - Only get correct update for last layer
 - Deep: compute derivative w.r.t. weight parameters that appear in the early layers but loss only computed at end
- Solution: error backpropagation (backprop)

BACKGROUND: COMPUTATION GRAPHS

- Represents process of computing mathematical expression broken down into separate operations
- Each operation is a node in the graph

COMPUTATION GRAPHS


- Example: $L(a,b,c) = c(a + 2b)$
 - $d = 2 * b$
 - $e = a + d$
 - $L = c * e$

COMPUTATION GRAPHS

- Example: $L(a,b,c) = c(a + 2b)$
 - $d = 2 * b$
 - $e = a + d$
 - $L = c * e$
- Represent steps as graph
 - Nodes = operations
 - Directed edges = show output from operation as input to next
- Forward pass: apply each operation left to right, passing outputs forward to next node

COMPUTATION GRAPHS

- Example: $L(a=3, b=1, c=-2) = c(a + 2b)$
 - $d = 2 * b$
 - $e = a + d$
 - $L = c * e$



BACKWARD DIFFERENTIATION ON COMPUTATION GRAPHS

- **Backward pass**: used to compute derivatives for weight update
- Example: compute derivative of output function L w.r.t. each input variable

BACKWARD DIFFERENTIATION ON COMPUTATION GRAPHS

- **Backward pass**: used to compute derivatives for weight update
- Example: compute derivative of output function L w.r.t. each input variable
 - Derivatives tell how much small change in variable affects L
 - $\frac{\partial L}{\partial a} \quad \frac{\partial L}{\partial b} \quad \frac{\partial L}{\partial c}$

BACKWARD DIFFERENTIATION ON COMPUTATION GRAPHS

- **Backward pass**: used to compute derivatives for weight update
- Example: compute derivative of output function L w.r.t. each input variable
 - Derivatives tell how much small change in variable affects L
 - $\frac{\partial L}{\partial a} \quad \frac{\partial L}{\partial b} \quad \frac{\partial L}{\partial c}$
 - Using $L = ce$ & chain rule: $\frac{\partial L}{\partial c} = e, \frac{\partial L}{\partial a} = \frac{\partial L}{\partial e} \frac{\partial e}{\partial a}, \frac{\partial L}{\partial b} = \frac{\partial L}{\partial e} \frac{\partial e}{\partial d} \frac{\partial d}{\partial b}$

BACKWARD DIFFERENTIATION ON COMPUTATION GRAPHS

- **Backward pass**: used to compute derivatives for weight update
- **Example**: compute derivative of output function L w.r.t. each input variable

- Derivatives tell how much small change in variable affects L

- $\frac{\partial L}{\partial a} \quad \frac{\partial L}{\partial b} \quad \frac{\partial L}{\partial c}$

- Using $L = ce$ & chain rule: $\frac{\partial L}{\partial c} = e$, $\frac{\partial L}{\partial a} = \frac{\partial L}{\partial e} \frac{\partial e}{\partial a}$, $\frac{\partial L}{\partial b} = \frac{\partial L}{\partial e} \frac{\partial e}{\partial d} \frac{\partial d}{\partial b}$

BACKWARD DIFFERENTIATION ON COMPUTATION GRAPHS

- Using $L = ce$ & chain rule: $\frac{\partial L}{\partial c} = e$, $\frac{\partial L}{\partial a} = \frac{\partial L}{\partial e} \frac{\partial e}{\partial a}$, $\frac{\partial L}{\partial b} = \frac{\partial L}{\partial e} \frac{\partial e}{\partial d} \frac{\partial d}{\partial b}$

$$\begin{aligned} L = ce : \quad & \frac{\partial L}{\partial e} = c, \frac{\partial L}{\partial c} = e \\ e = a + d : \quad & \frac{\partial e}{\partial a} = 1, \frac{\partial e}{\partial d} = 1 \\ d = 2b : \quad & \frac{\partial d}{\partial b} = 2 \end{aligned}$$

BACKWARD DIFFERENTIATION ON COMPUTATION GRAPHS


- Using $L = ce$ & chain rule: $\frac{\partial L}{\partial c} = e$, $\frac{\partial L}{\partial a} = \frac{\partial L}{\partial e} \frac{\partial e}{\partial a}$, $\frac{\partial L}{\partial b} = \frac{\partial L}{\partial e} \frac{\partial e}{\partial d} \frac{\partial d}{\partial b}$

$$\begin{aligned} L = ce : \quad & \frac{\partial L}{\partial e} = c, \frac{\partial L}{\partial c} = e \\ e = a + d : \quad & \frac{\partial e}{\partial a} = 1, \frac{\partial e}{\partial d} = 1 \\ d = 2b : \quad & \frac{\partial d}{\partial b} = 2 \end{aligned}$$

- To compute backward pass
 - Compute each partial along each edge of CG from right to left
 - Multiply necessary partials to get final derivative needed

BACKWARD DIFFERENTIATION ON COMPUTATION GRAPHS


$$\begin{aligned}
 L = ce : \quad & \frac{\partial L}{\partial e} = c, \frac{\partial L}{\partial c} = e \\
 e = a + d : \quad & \frac{\partial e}{\partial a} = 1, \frac{\partial e}{\partial d} = 1 \\
 d = 2b : \quad & \frac{\partial d}{\partial b} = 2
 \end{aligned}$$



BACKWARD DIFFERENTIATION ON COMPUTATION GRAPHS

$$\begin{aligned}
 L = ce : \quad & \frac{\partial L}{\partial e} = c, \frac{\partial L}{\partial c} = e \\
 e = a + d : \quad & \frac{\partial e}{\partial a} = 1, \frac{\partial e}{\partial d} = 1 \\
 d = 2b : \quad & \frac{\partial d}{\partial b} = 2
 \end{aligned}$$

1. Compute local partial derivative w.r.t. parent
2. Multiply it by partial derivative passed down from parent
3. Pass value to child



BACKWARD DIFFERENTIATION FOR NN

BACKWARD DIFFERENTIATION FOR NN

- Binary classification (sigmoid)


$$z^{[1]} = W^{[1]}x + b^{[1]}$$

$$a^{[1]} = \text{ReLU}(z^{[1]})$$

$$z^{[2]} = W^{[2]}a^{[1]} + b^{[2]}$$

$$a^{[2]} = \sigma(z^{[2]})$$

$$\hat{y} = a^{[2]}$$




BACKWARD DIFFERENTIATION FOR NN

- Activation function derivatives

$$\frac{d\sigma(z)}{dz} = \sigma(z)(1 - \sigma(z))$$

$$\frac{d \tanh(z)}{dz} = 1 - \tanh^2(z)$$

$$\frac{d \text{ReLU}(z)}{dz} = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$$



NN LEARNING

- NN Optimization is non-convex & more complex than LR
- Initialize weights with small random numbers
- Normalize input values to have 0 mean & unit variance
- Use **dropout** regularization to help avoid overfitting
 - Randomly drop some units & their connections from the network during training
- Tune **hyperparameters** on devset
 - NN parameters = W & b learned by GD
 - Learning rate η , mini-batch size, architecture, regularization ...
 - # of layers, # of hidden nodes per layer, activation functions

- Neural Units
 - Biological neuron → artificial neuron
 - Inputs (x, w, b)
 - Activation functions
 - Sigmoid, tanh, ReLU
- XOR Problem
 - Unit → networks
- Feed forward NN
 - No cycles
 - Hidden layer & matrix computation
 - Output calculation
- Training NN
 - Cross-entropy loss, gradient descent, back prop

SUMMARY

Units

XOR Problem

Feed Forward Neural Network

Training NN

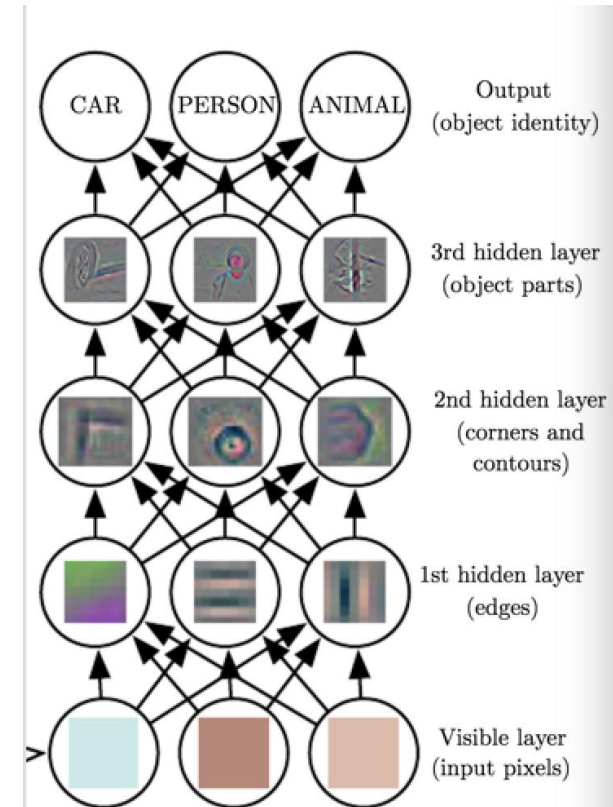
RNN & LSTM

NN

SIMPLE RNN

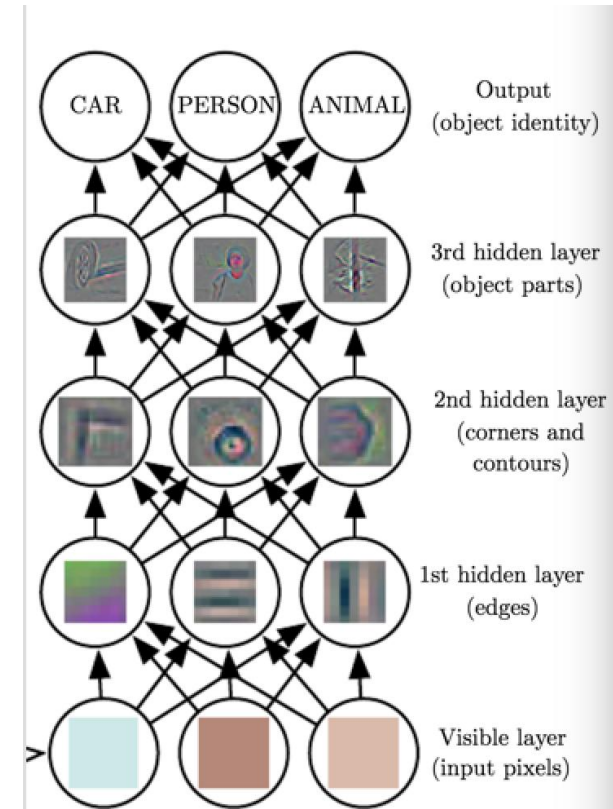
- RNN: any network that contains a **cycle** within its network connections
- Cycle: output value of unit/node is in/directly dependent on earlier outputs (as its input)
- **Elman** (1990) or **simple recurrent networks**
 - Effective for spoken & written language
 - Base for Encoder-Decoder models & QA models

FEEDFORWARD RECAP




FEEDFORWARD RECAP

- Training
 - Input units represent info
 - Multiply by weights
 - If the sum of weights $>$ threshold (activates) triggers next units




FEEDFORWARD RECAP

- **Training**
 - Input units represent info
 - Multiply by weights
 - If the sum of weights $>$ threshold (activates) triggers next units
- **Learning (Backprop)**
 - Compare the output network produces (y^*) with output should have produced (y)
 - Use difference between them to modify weights (work backwards)




SIMPLE RNN



SIMPLE RNN

current input's features * weight matrix




SIMPLE RNN

pass through activation function to
compute activation value a for h



current input's features * weight matrix



SIMPLE RNN


h calculates output value



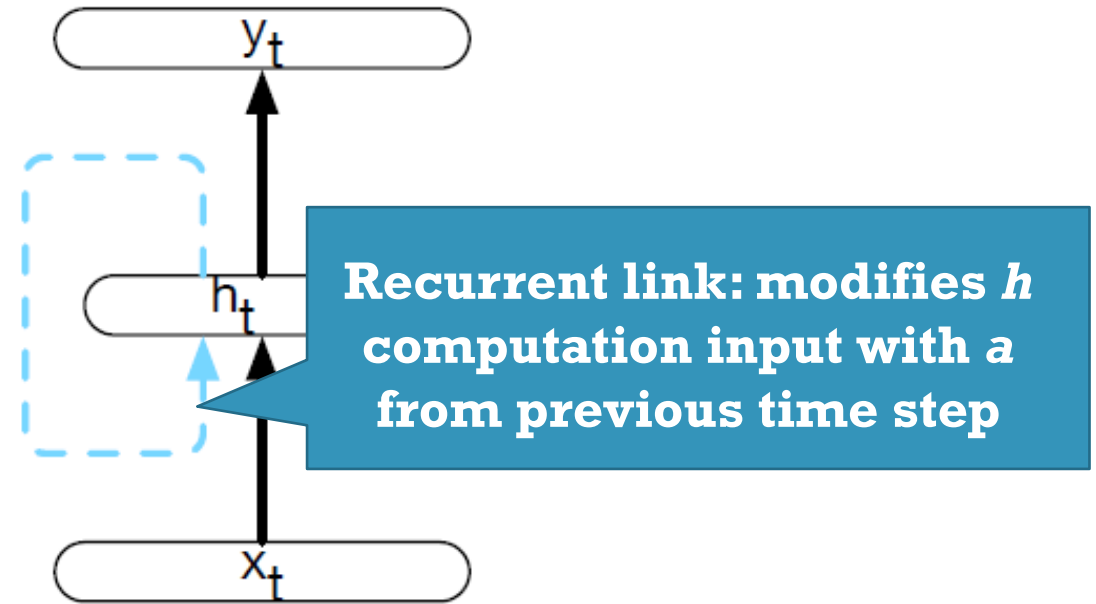
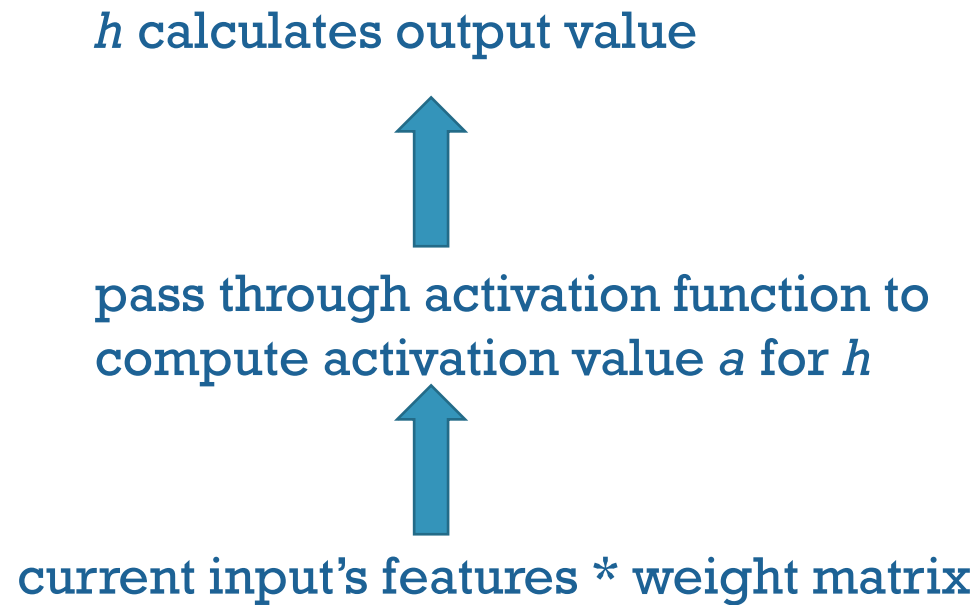
pass through activation function to
compute activation value a for h



current input's features * weight matrix



SIMPLE RNN




RECURRENT LINK

- Previous timestep's hidden layer \approx memory, context
- Encodes earlier processing steps
- Helps make future decisions
- Key: doesn't limit length of prior context
 - Context “remembered” in the previous hidden layer includes info all the way back to the beginning of the sequence


SIMPLE RNN

- Given
 - Input vector (x_t)
 - Values for h from previous time step
- Still standard feedforward



SIMPLE RNN

- Given
 - Input vector (x_t)
 - Values for h from previous time step
- Still standard feedforward
- Key change: new set of weights U
 - Connects previous timestep hidden layer (h_{t-1}) to current hidden layer (h_t)
 - Weights U determine how network should use past context to calculate output for current input




INFERENCE

- Forward inference: map sequence of inputs to sequence of outputs
- Similar to feedforward network
- Output $y_t = \text{input } x_t + \text{activation value for hidden layer } h_t$

INFERENCE

- $h_t = g\{(x_t * W) + (h_{t-1} * U)\}$
- $y_t = \text{softmax}(h_t * V)$



INFERENCE

- $h_t = g\{(x_t * W) + (h_{t-1} * U)\}$
- $y_t = \text{softmax}(h_t * V)$

function FORWARDRNN($x, network$) **returns** output sequence y


$h_0 \leftarrow 0$

for $i \leftarrow 1$ **to** LENGTH(x) **do**


$h_i \leftarrow g(U h_{i-1} + W x_i)$

$y_i \leftarrow f(V h_i)$


return y




INFERENCE



INFERENCE



INFERENCE



TRAINING

- Just like feedforward networks
 - Training set
 - Loss function
 - Backprop

TRAINING

- Just like feedforward networks
 - Training set
 - Loss function
 - Backprop
- Weights
 - W : from input layer to hidden layer
 - U : from previous hidden layer to current hidden layer
 - V : from hidden layer to output layer

BACKPROP

- 2 new concerns
 - (1) To calculate loss for output at time t we need hidden layer from $t-1$


BACKPROP

- 2 new concerns
 - (1) To calculate loss for output at time t we need hidden layer from $t-1$
 - (2) Hidden layer at time t influences output at time t & hidden layer at time $t+1$

BACKPROP


- 2 new concerns
 - (1) To calculate loss for output at time t we need hidden layer from $t-1$
 - (2) Hidden layer at time t influences output at time t & hidden layer at time $t+1$
 - So it also influences loss at $t+1$
 - To calculate error accruing in h_t we need to know influence on current output & those that follow

BACKPROP




BACKPROP THROUGH TIME

- Input/output example pair at $t = 2$
- What do we need to compute gradients to update U , V , & W ?



V UPDATE


- Gradients to update V
- Same as in FFN
- Derivative of loss w.r.t. weights V



V UPDATE

- Gradients to update V
- Same as in FFN
- Derivative of loss w.r.t. weights V

$$\boxed{\frac{\partial L}{\partial V}} = \frac{\partial L}{\partial a} \frac{\partial a}{\partial z} \frac{\partial z}{\partial V}$$




V UPDATE

- Gradients to update V
- Same as in FFN
- Derivative of loss w.r.t. weights V

$$\frac{\partial L}{\partial V} = \frac{\partial L}{\partial a} \frac{\partial a}{\partial z} \frac{\partial z}{\partial V}$$

Derivative of loss function
w.r.t network output a




V UPDATE

- Gradients to update V
- Same as in FFN
- Derivative of loss w.r.t. weights V

$$\frac{\partial L}{\partial V} = \frac{\partial L}{\partial a} \frac{\partial a}{\partial z} \frac{\partial z}{\partial V}$$

Derivative of network output a
w.r.t. intermediate activation z




V UPDATE

- Gradients to update V
- Same as in FFN
- Derivative of loss w.r.t. weights V

$$\frac{\partial L}{\partial V} = \frac{\partial L}{\partial a} \frac{\partial a}{\partial z} \frac{\partial z}{\partial V}$$

Derivative of intermediate activation w.r.t weights V



V UPDATE

- δ_{out} : error term that represents how much loss is associated with each of the units in output layer

$$\frac{\partial L}{\partial V} = \left(\frac{\partial L}{\partial a} \frac{\partial a}{\partial z} \right) \frac{\partial z}{\partial V}$$

- Final gradient needed to update V

$$\frac{\partial L}{\partial V} = \delta_{out} h_t$$

U & W UPDATE

- Difference from FFN: computing W & U
- h_t contributes to output & error at time t & $t+1$
- So δ_h must include error from both timesteps

$$\delta_h = g'(z)V\delta_{out} + \delta_{next}$$

U & W UPDATE

- Difference from FFN: computing W & U
- h_t contributes to output & error at time t & $t+1$
- So δ_h must include error from both timesteps

$$\frac{\partial L}{\partial W} = \delta_h x_t$$
$$\frac{\partial L}{\partial U} = \delta_h h_{t-1}$$

U & W UPDATE

- Difference from FFN: computing W & U
- h_t contributes to output & error at time t & $t+1$
- So δ_h must include error from both timesteps

$$\frac{\partial L}{\partial W} = \delta_h x_t$$
$$\frac{\partial L}{\partial U} = \delta_h h_{t-1}$$

- Compute error: “assign proportional blame”
 - Backprop δ_h to previous h_{t-1}
 - Proportional based on U

$$\delta_{next} = g'(z)U\delta_h$$

TRAINING SUMMARY

- Backpropagation Through Time
- First pass
 - Do forward inference: compute h_t & y_t
 - Accumulate loss at each step
 - Save value of h_t at each step to use in next timestep

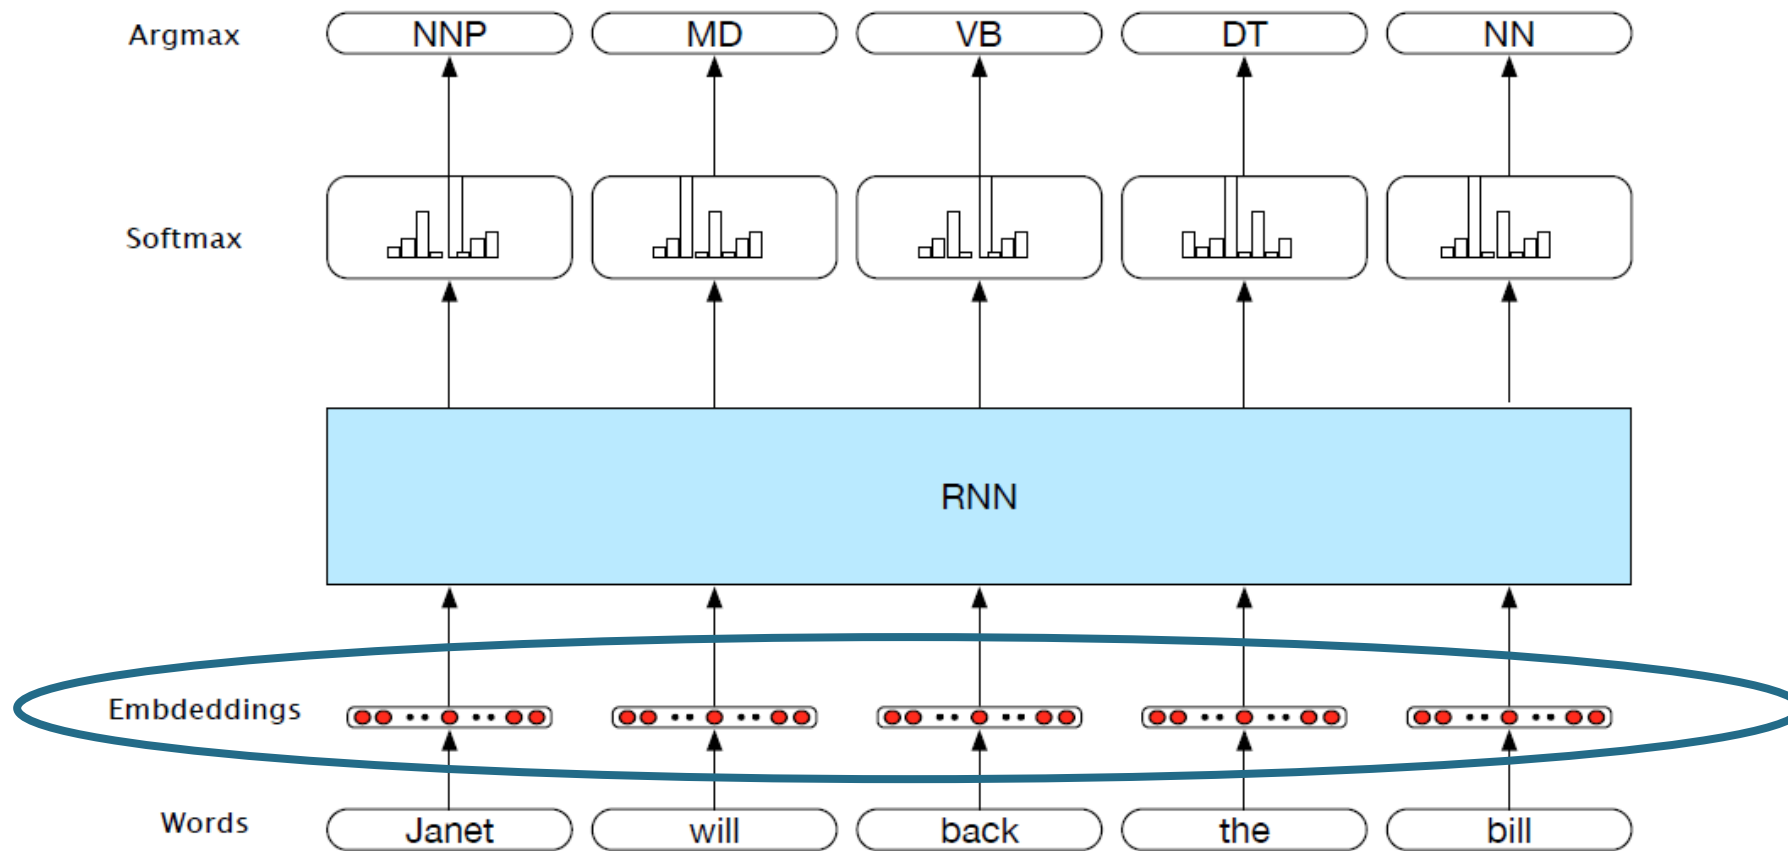
TRAINING SUMMARY

- Backpropagation Through Time
 - First pass
 - Do forward inference: compute h_t & y_t
 - Accumulate loss at each step in time
 - Save value of h_t at each step to use in next timestep
 - Second pass
 - Process sequence in reverse
 - Compute required error term gradients
 - Compute & save error term for use in hidden layer for each backward step in time

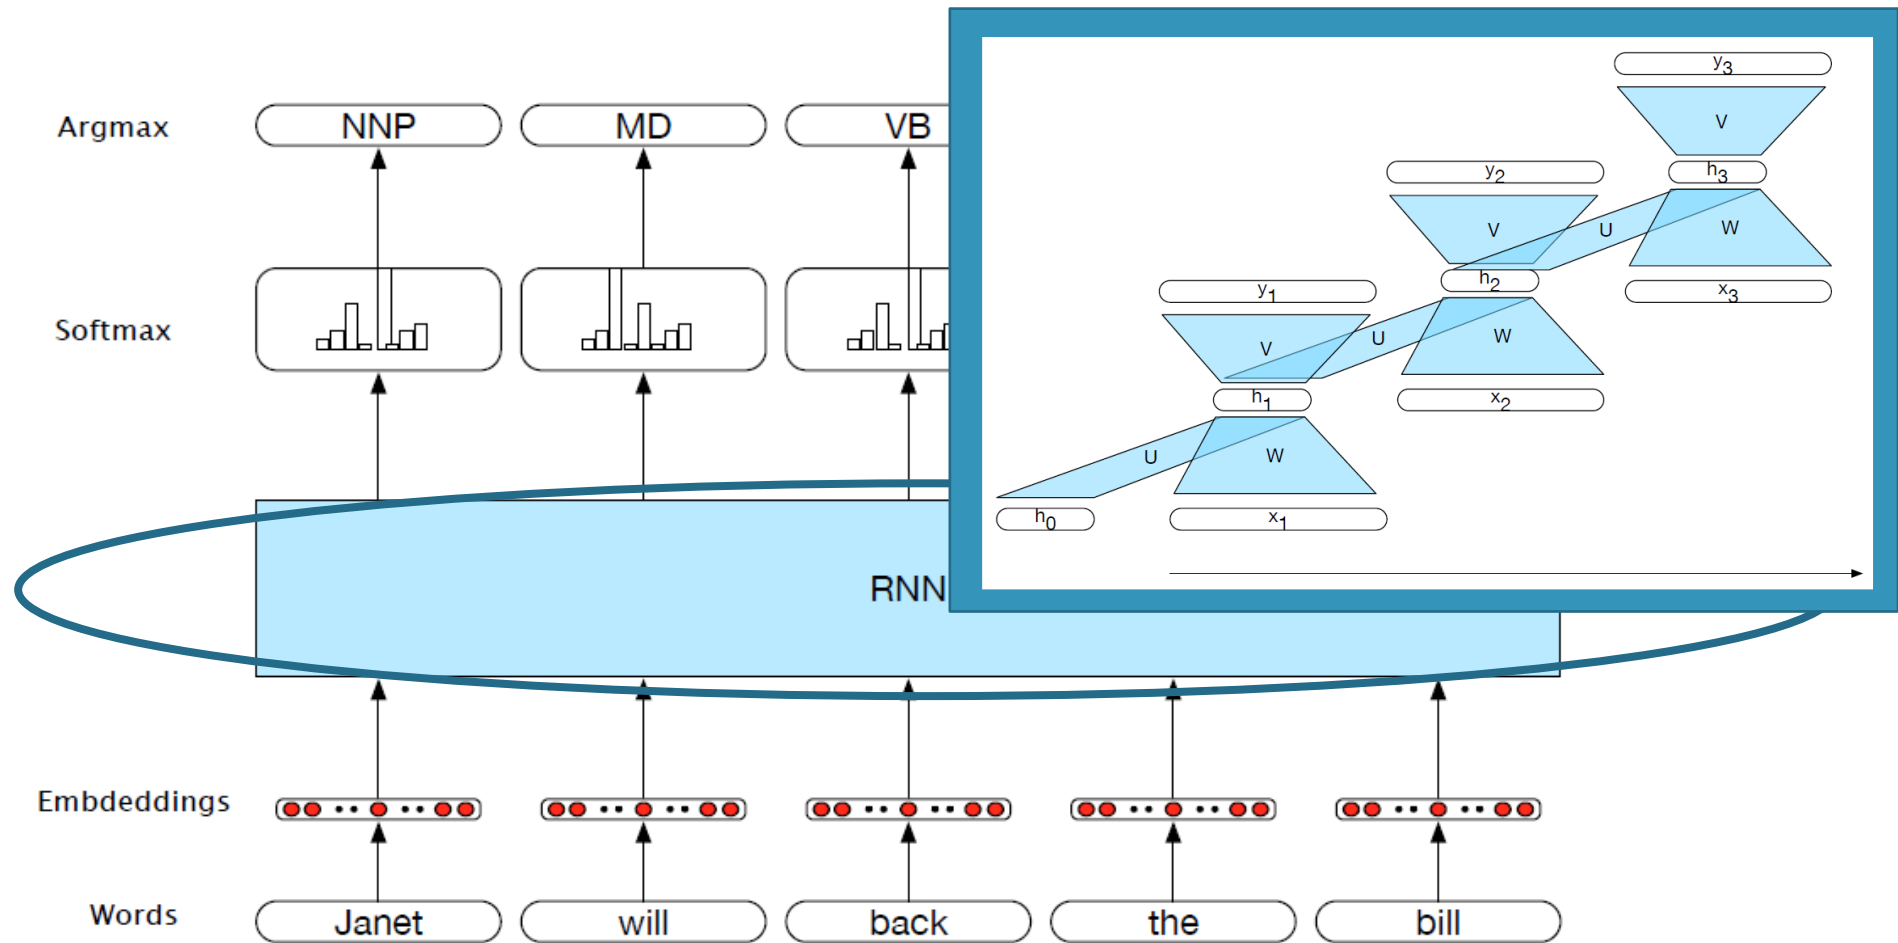
RNN APPLICATIONS

- Effective for
 - Language modeling
 - Sequence labeling tasks (e.g., POS tagging)
 - Sequence classification tasks (sentiment analysis, topic classification)
- Basis for sequence-to-sequence approaches
 - Summarization
 - Machine Translation
 - Question Answering

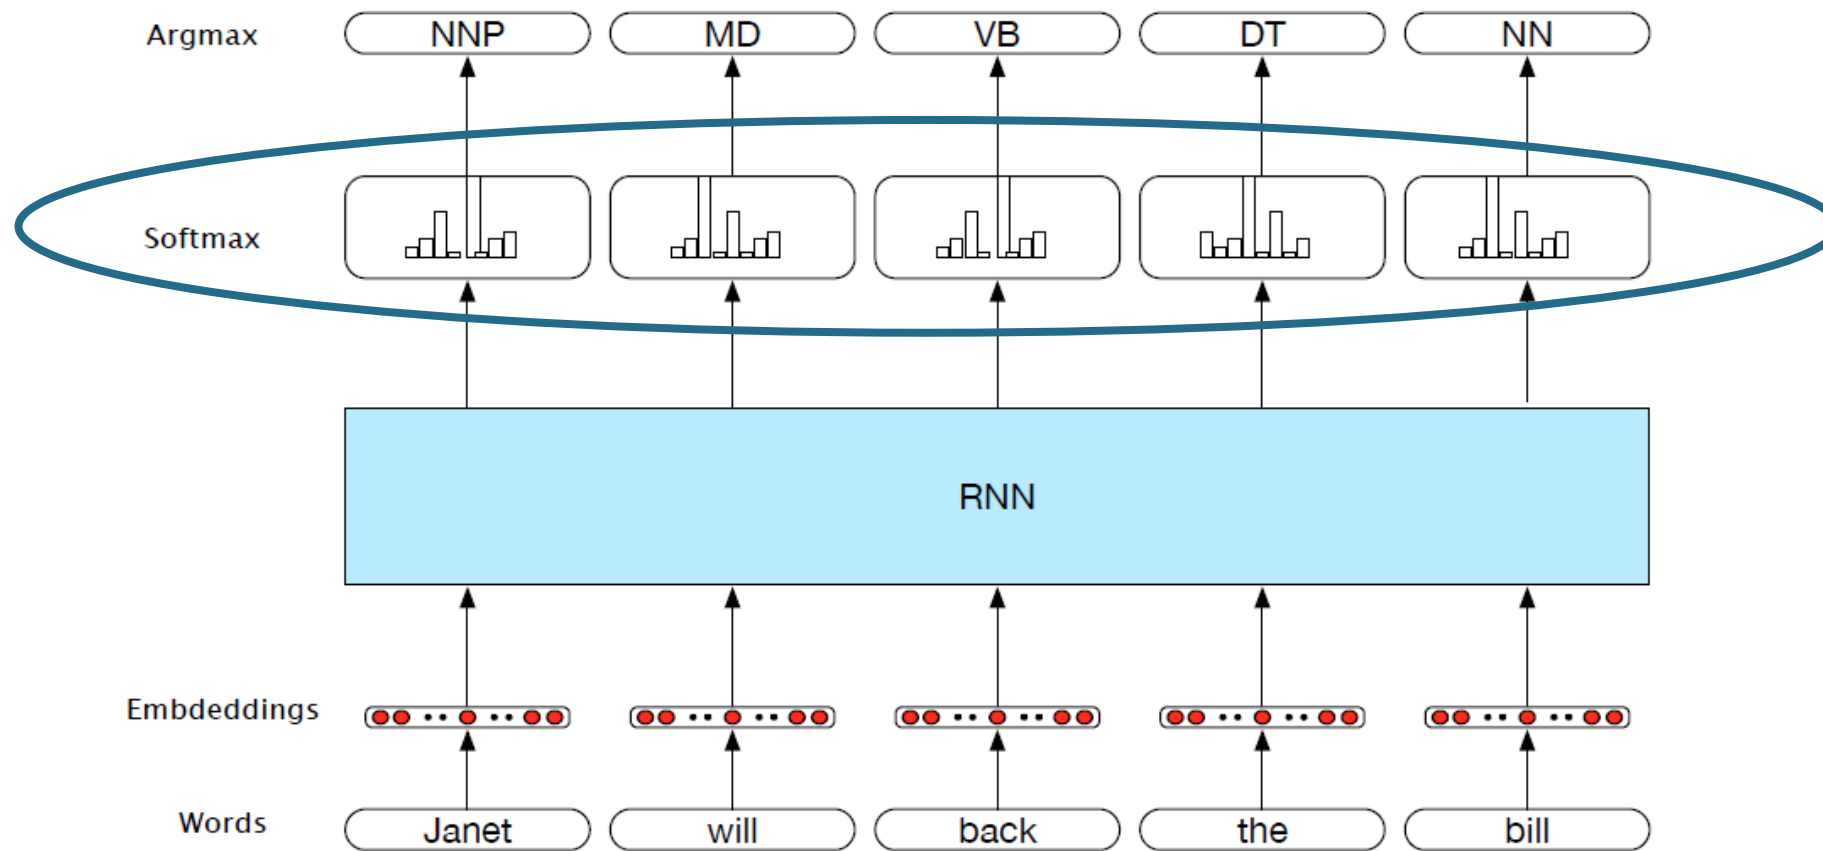
SEQUENCE LABELING



SEQUENCE LABELING

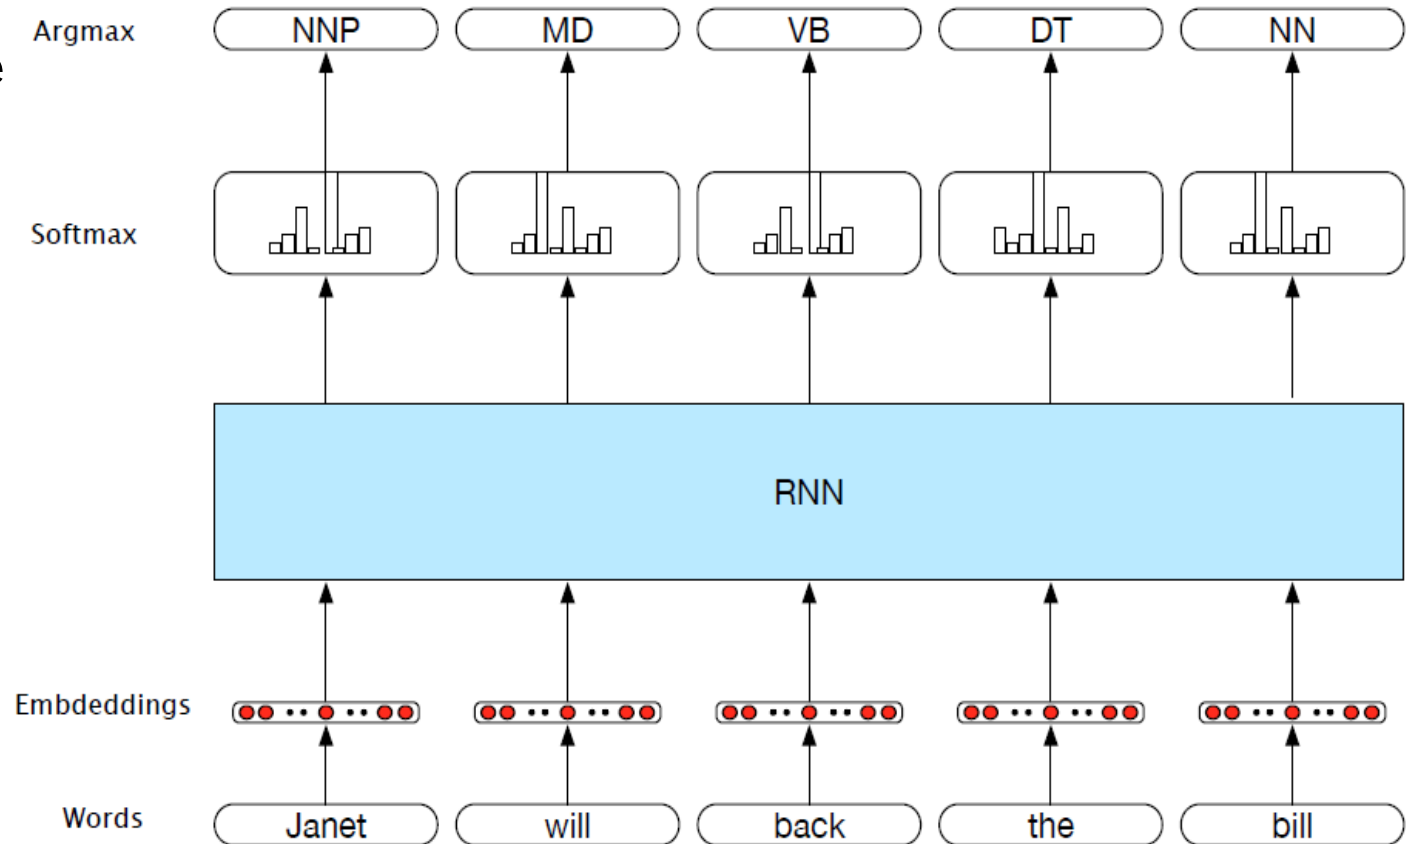


SEQUENCE LABELING




SEQUENCE LABELING

- To generate a tag sequence for a given input
 - Run forward inference over input sequence
 - Select most likely tag from *softmax* at each step



DEEP NEURAL NETWORK

- **Deep NN** = simple RNN with feedforward classifier
 - Stacked RNNs
 - Bidirectional RNNs
- **End-to-end training**: uses loss from downstream apps to adjust weights all the way throughout the network

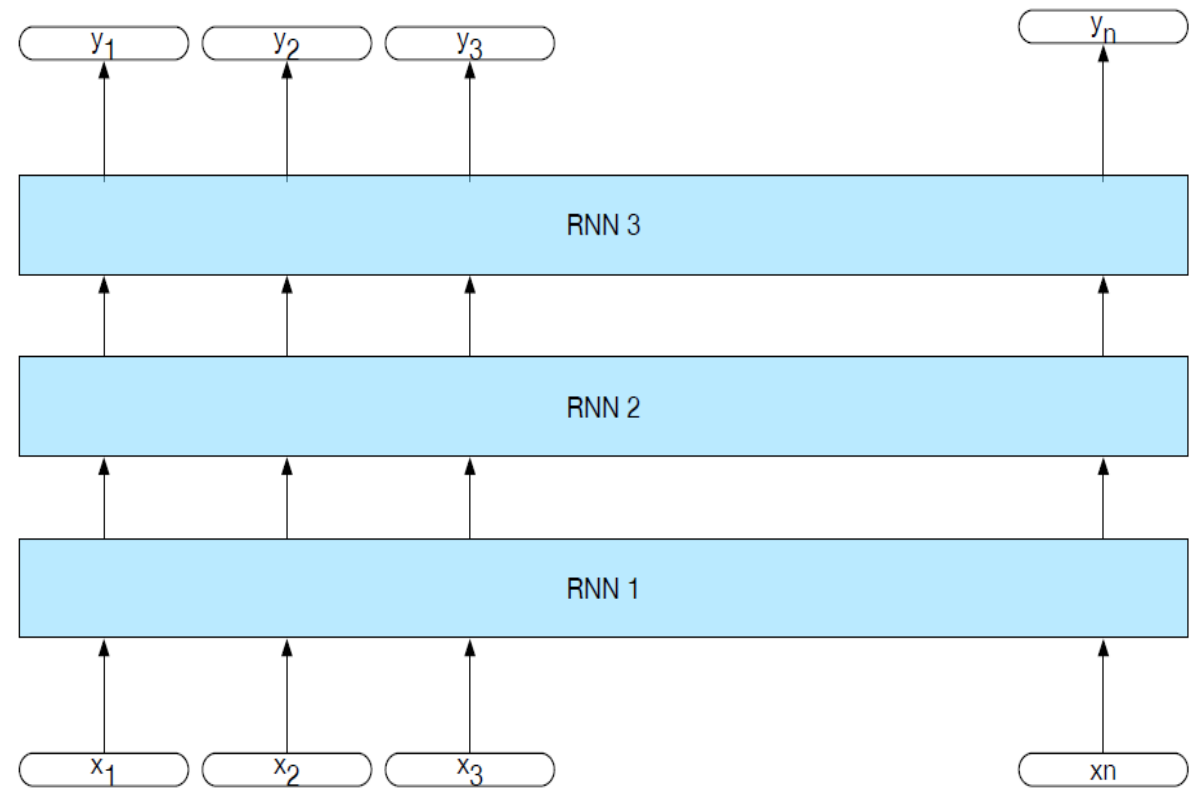


DEEP NEURAL NETWORK

- No intermediate outputs for words in sequence preceding last element → no loss terms associated with those
- Loss function based entirely on final text classification task
 - *Softmax* output (from FFN) + cross-entropy loss → training
 - Classification error is backpropagated through all aspects of FFN: weights in FF classifier → input → RNN 3 matrices (U, V, W)

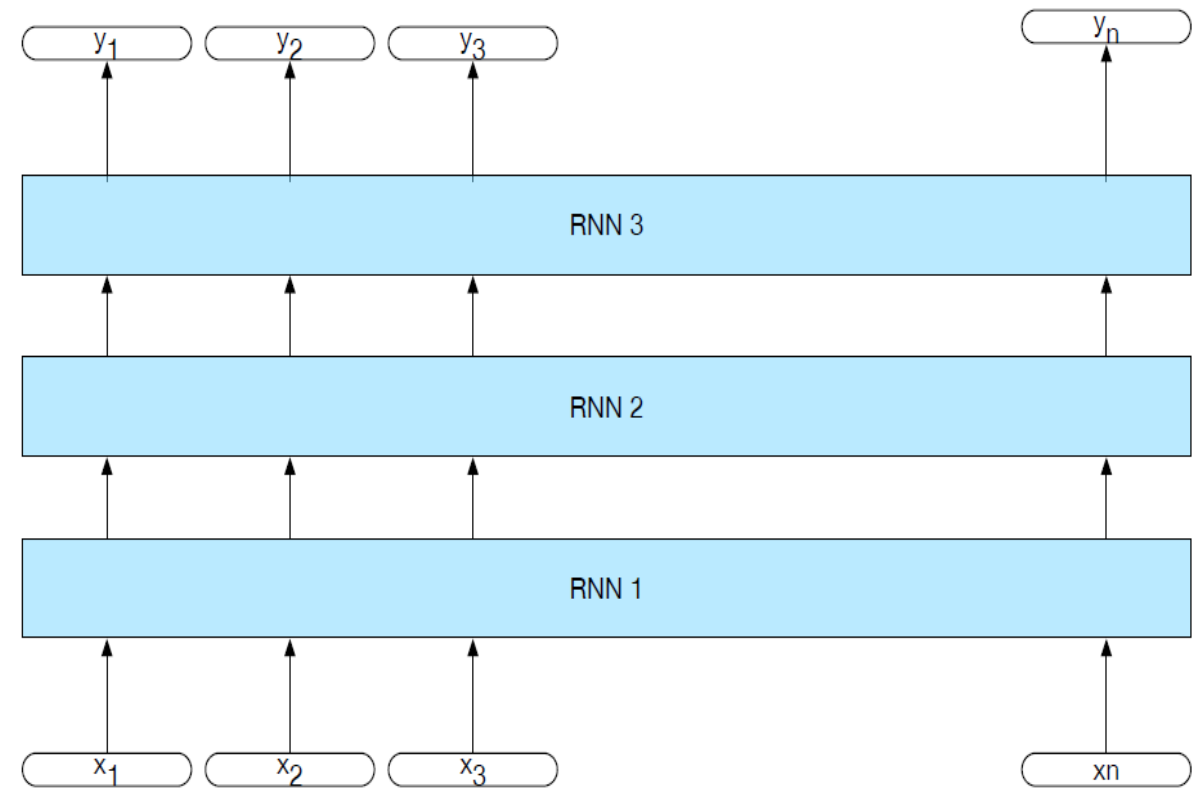
STACKED RNN

- Up until now:
 - RNN input = word sequences or embeddings (vectors)
 - RNN output = vectors for predicting words, tags, sequence labels
- Why not use entire sequence of outputs from 1 RNN as input sequence to another?



STACKED RNN

- Up until now:
 - RNN input = word sequences or embeddings (vectors)
 - RNN output = vectors for predicting words, tags, sequence labels
- Why not use entire sequence of outputs from 1 RNN as input sequence to another?
- **Stacked RNN**: multiple networks where output layer of 1 layer serves as input to subsequent layer
 - # of stacks task & training set specific
 - # of stacks rises, training costs rises
 - Induces representations at differing levels of abstractions across layers



BIDIRECTIONAL RNN (BI-RNN)

- In simple RNN hidden state at time t represents everything network knows about sequence up to that point
- Think of it as context to the **left** of the current time

$$h_t^f = RNN_{forward}(x_1^t)$$

BIDIRECTIONAL RNN (BI-RNN)

- In simple RNN hidden state at time t represents everything network knows about sequence up to that point
- Think of it as context to the **left** of the current time

$$h_t^f = RNN_{forward}(x_1^t)$$

- If we have access to entire input sequence at once, use context to the right too
- To grab it, we train an RNN on input sequence *in reverse*

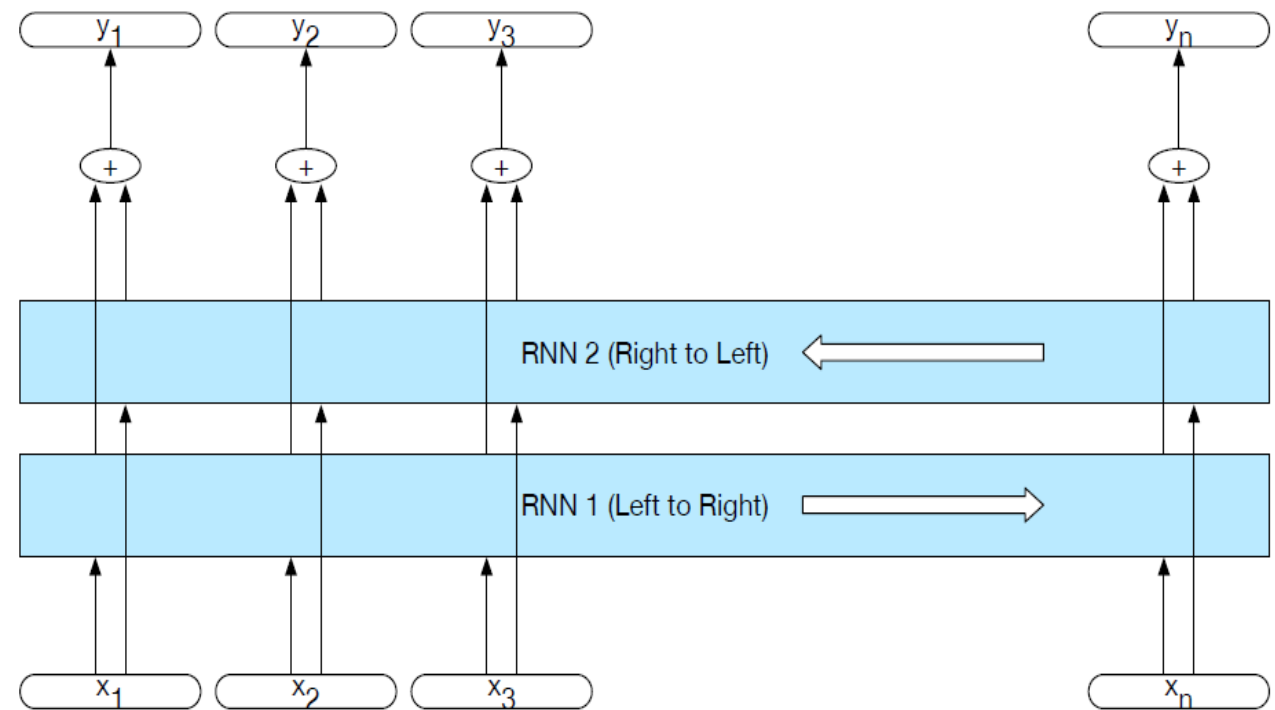
$$h_t^b = RNN_{backward}(x_t^n)$$

BI-RNN

- Bidirectional RNN = forward information + backward
- 2 independent RNNs
 - Input processed start to end
 - Input processed end to start
- Output combined into single representation that captures left & right contexts of input at each point in time


BI-RNN

- Outputs of forward & backward pass are concatenated
- Output at each time step captures info to the left & right of current input
- Example: use concatenated output as basis for local labeling decision in sequence labeling apps



BI-RNN

- Highly effective for sequence classification
- Final state naturally reflects more info about end of sentence than its beginning
 - Previous attempt: input final RNN's hidden state to FF classifier
- Bi-RNN solution: combine final forward & backward hidden states & use as input



MANAGING CONTEXT

- Difficult to train RNNs for tasks that need information far away from current position in processing
- Have access to entire preceding sequence
- But info encoded in hidden states is usually *local* i.e., more relevant to most recent parts of input sequence & recent decisions

MANAGING CONTEXT

- Usually though distant information is important for NLP applications
- LM example: *The flights the airline was cancelling were full.*
 - $P(\text{was} \mid \dots \text{airline})$: makes sense because verb matches
 - $P(\text{were} \mid \dots \text{airline})$: trickier because *flights* further away & singular *airline* is closer
- Ideally network should be able to retain distant info until needed while processing intermediate parts of sequence correctly

MANAGING CONTEXT

- RNNs have trouble carrying distant information forward
 - Hidden layers (& weights that determine their value) are asked to handle 2 tasks simultaneously
 - Provide useful info for current decision
 - Update & carry forward info for future decisions

MANAGING CONTEXT

- RNNs have trouble carrying distant information forward
 - Hidden layers (& weights that determine their value) are asked to handle 2 tasks simultaneously
 - Provide useful info for current decision
 - Update & carry forward info for future decisions
- Need to backprop error through time
 - h_t contributes to loss at $t+l$
 - During backward pass, hidden layers are multiplied repeatedly
 - Result: **vanishing gradient problem**

MANAGING CONTEXT

- RNNs have trouble carrying distant information forward
 - Hidden layers (& weights that determine their value) are asked to handle 2 tasks simultaneously
 - Provide useful info for current decision
 - Update & carry forward info for future decisions
 - Need to backprop error through time
 - h_t contributes to loss at $t+l$
 - During backward pass, hidden layers are multiplied repeatedly
 - Result: **vanishing gradient problem**
- Need a way to forget info we don't need anymore & remember info we'll need in the future

LONG SHORT-TERM MEMORY (LSTM)

- Divide context management problem into 2 subproblems
 - Remove info that's no longer needed
 - Add info likely to be needed for later decision making
- Key to solving both: learn how to manage context instead of hard-coding it into architecture

LONG SHORT-TERM MEMORY (LSTM)

- (1) Add explicit context layer to architecture
- (2) Use specialized neural units with **gates** to control info flow through the units in layers
 - Implemented through additional weights that operate sequentially on input & previous hidden & context layers

LSTM GATES

- Design
 - Feedforward layer
 - Sigmoid activation function
 - Pointwise multiplication with layer being gated
- Forget gate
- Add gate
- Output gate

LSTM GATES

- Design
 - Feedforward layer
 - Sigmoid activation function
 - Pointwise multiplication with layer being gated
- Forget gate
- Add gate
- Output gate

**Sigmoid because it
pushes output to 0 or 1**

LSTM GATES

- Design
 - Feedforward layer
 - Sigmoid activation function
 - **Pointwise multiplication** with layer being gated
- Forget gate
- Add gate
- Output gate

Sigmoid + PM \approx binary mask
Values align near 1 pass; lower erased

LSTM GATES

- Forget gate: Deletes info from context that's no longer needed
 - Computes weighted sum of h_{t-1} + current input
 - Passes that value through sigmoid \rightarrow mask
 - Multiply that mask by context vector \rightarrow removes info
- Computation step
- Add gate
- Output gate

LSTM GATES

- Forget gate: Deletes info from context that's no longer needed
- Computation step: computes info needed from previous hidden state & current inputs using *tanh*
- Add gate
- Output gate

LSTM GATES

- Forget gate: Deletes info from context that's no longer needed
- Computation step: computes info needed
- Add gate: selects information to add to current context
 - Computes weighted sum of previous hidden layer & current input
 - Passes that value through sigmoid → mask
 - Adds mask to context vector → new context vector with new info
- Output gate

LSTM GATES

- Forget gate: Deletes info from context that's no longer needed
- Computation step: computes info needed
- Add gate: selects information to add to current context
- Output gate: decides what info is needed for current hidden state

GATED RECURRENT UNIT (GRU)

- LSTM requires learning 8 weights
 - U & W for each of the 4 gates within each unit
- GRUs
 - Drop separate context vector
 - Reduce number of gates to 2
 - Reset gate (r)
 - Update gate (z)

GRU GATES

- Like LSTM: uses sigmoid to create binary-like mask
 - Blocks info if values near zero
 - Allows info to pass through unchanged if values near 1


GRU GATES

- Like LSTM: uses sigmoid to create binary-like mask
 - Blocks info if values near zero
 - Allows info to pass through unchanged if values near 1
- Reset gate
 - Decides which aspects of previous hidden state are relevant to current context (or should be ignored)
 - Computes mask to get intermediate new hidden state

GRU GATES


- Like LSTM: uses sigmoid to create binary-like mask
 - Blocks info if values near zero
 - Allows info to pass through unchanged if values near 1
- Reset gate
 - Decides which aspects of previous hidden state are relevant to current context (or should be ignored)
 - Computes mask to get intermediate new hidden state
- Update gate
 - Decides which aspects of new state will be used directly in new hidden state
 - Decides which aspects of previous state are preserved for future use

MODULARITY




(a)

Feedforward




(b)

Simple RNN




(c)



(d)


MODULARITY



(a)


Feedforward

MODULARITY



(a)

Feedforward




(b)


Simple RNN

MODULARITY

- Complexity encapsulated within neural unit
- Modularity allows LSTM & GRU units to be used in other network architectures
- Multi-layer networks using gated units can be unrolled into deep feedforward networks



(c)



(d)

- Simple RNN
 - Inference
 - Training
- Applications
 - RNLM
 - Sequence labeling
 - POS tagging
 - NER
 - Sequence Classification
- Deep Networks
 - Stacked RNNs
 - Bidirectional RNNs
- Managing Context
 - LSTMs
 - GRUs

SUMMARY